

Team KASS Project Reflection

1. Please list out changes in the directions of your project if the final project is different from your original proposal (based on your stage 1 proposal submission).

There have been no changes to our project between our stage 1 proposal submission and our final project submission. We kept our final submission identical to the original project proposal since our original plan accounted for any potential pitfalls, so we were able to avoid having to make drastic changes throughout the semester.

One thing we proposed in Stage 1, which we did not add to our final project, was <https://www.kaggle.com/datasets/llui85/tokyo-2021-olympics-complete-grouped-by-type> was another dataset that we wanted to integrate onto our 10000+ row TA-approved dataset. We decided against this as the data was incompatible with our dataset, and it did not add much functionality. The joins etc. would have been too complicated due to the different data types (CSV vs. JSON) and many more problems, which all would have added to unnecessary problems.

We also did not implement this additional feature as it would be beyond the scope as the dataset did not have booking or event spaces. We did not have this data on the TA-approved dataset, but we did find data from another source. We added filtering and ticket count in our App.

Another thing we wanted to do was role-based authentication to do CRUD operations, but this was again an add-on feature. Instead, we directly do Authentication based access to perform CRUD operations.

Again, note that all the changes that were not implemented above were categorized as add-ons in our application. We delivered all the core functionality and added more features, such as visualization tools.

2. Discuss what you think your application achieved or failed to achieve regarding its usefulness.

Our application has successfully created a place where a user can sort through and analyze the 2020 Tokyo Olympics data. All while being able to add and delete athletes to the data set, update the medal counts of different countries, and filter through different events. Additionally, the visualizations provide easy-to-understand readability to the large sets of data in our application's athlete and medal sections.

We have succeeded on multiple fronts, and we do this in the following ways - Data and Functionality.

Data

We have different kinds of data, workflows, and access controls

- We cleaned and set up data for our application
- User data for authentication, authorization, and access control. It would include attributes such as username and password.

Functions

Any user has the following abilities that they can leverage from our website.

- One can log in to our website, and we process authentication and authorization.
- Anyone can leverage the search sort functionality of our website.
- Anyone can leverage the visualizations on the website.
- Anyone can request that certain athletes' data be added, updated, or deleted after authentication.
- Data enforcement for certain fields, such as countries, so random Non-existent countries can't be added.
- Have Stored Procedures and triggers to update stats and add functionality via UI, and we were able to visualize through the same.

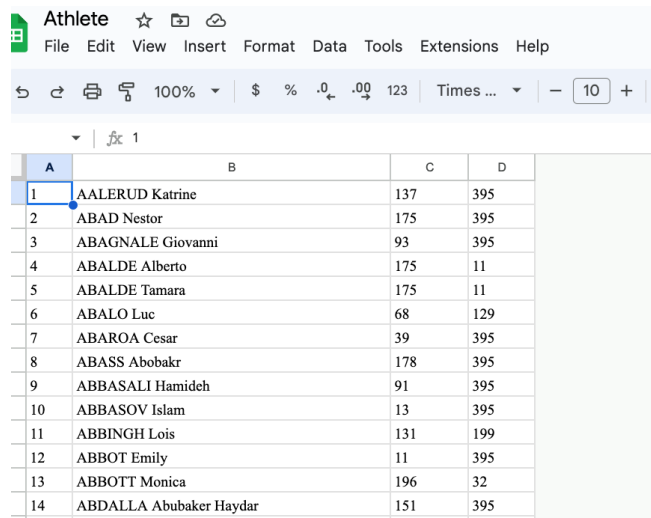
We set out to create an ultimate best-case website that functions better than the official website, and we were able to create a wonderful user experience for the same.

3. Discuss if you changed the schema or source of the data for your application

We had to modify the TA-approved dataset by introducing IDs to ensure that the tables have meaning between themselves, and this way, we were able to import the dataset directly into our backend and did not have to worry about complex queries and indexing concerns.

	A	B	C	D	E	F	G	H	I	J
1										
2										
3	Athletes	Coaches	Entries	Medals	Teams					
4	Name	Name	Discipline/Sport	Rank	Country					
5	Country	Country	Female Count	Gold Count	Discipline/Sport					
6	Discipline/Sport	Discipline/Sport	Male Count	Gold Count	Country					
7			Event	Total	Silver Count	Event				
8					Bronze Count					
9					Total Count					
10					Rank By Total					
11										
12										
13										
14	Login		Country		Sport	Athletes	Coaches			
15	Username		Name		Name	Name	Name			
16	Password		Sport		Medal Count	Male/Female	Country			
17			Athlete		Sport	Discipline/Sport				
18	Name		Medals		Medal Count	Country	Event			
19	Email				Team/Individual	Country	Athlete			
20										
21										
22										
23										
24										
25										
26										
27										
28	Login			Coach						
29										
30										
31										
32										
33										
34	Athlete									
35										
36										
37										

Above is the original schema, and we have the data to be mapped only via FKs across wherever applicable



A	B	C	D
1	AALERUD Katrine	137	395
2	ABAD Nestor	175	395
3	ABAGNALE Giovanni	93	395
4	ABALDE Alberto	175	11
5	ABALDE Tamara	175	11
6	ABALO Luc	68	129
7	ABAROA Cesar	39	395
8	ABASS Abobakr	178	395
9	ABBASALI Hamideh	91	395
10	ABBASOV Islam	13	395
11	ABBINGH Lois	131	199
12	ABBOT Emily	11	395
13	ABBOTT Monica	196	32
14	ABDALLA Abubaker Haydar	151	395

But this can be categorized as cleaning, and the original data was kept intact. However, we had to generate data for User tables synthetically. Additionally, we had to add a new dataset for the event table as it did not come under the original table or dataset.

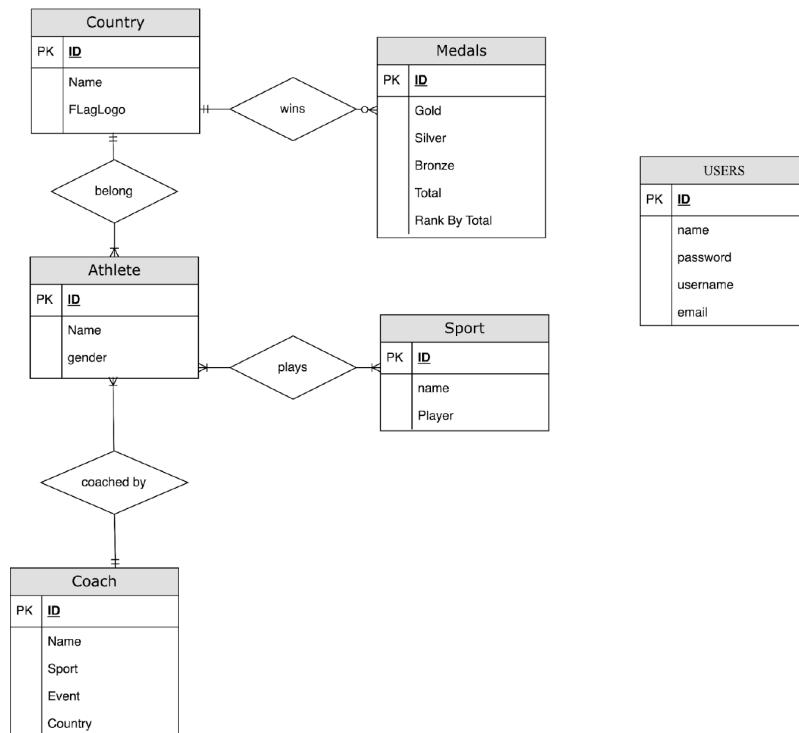
4. Discuss what you change to your ER diagram and/or your table implementations. What are some differences between the original design and the final design? Why? What do you think is a more suitable design?

We spent a lot of time designing the ER diagram perfectly the first time, and we discussed a lot about it. Our ER diagram is not changed from the proposal and is maintained intact as is. By making sure that we spent a lot of time designing the proper schema, we could design the entire schema and app properly the first time without any modifications. The DDL statements also remained intact. The only change in the databases is that we had to modify the athlete table once

```
CREATE TABLE Athlete (
  ID INT PRIMARY KEY,
  NAME VARCHAR(100) ,
  COUNTRYID INT,
  COACHID INT,
  FOREIGN KEY (COUNTRYID) REFERENCES Country(ID),
  FOREIGN KEY (COACHID) REFERENCES Coach(ID)
);
```

As ID did not have auto-increment and we had to alter the table to add AUTO_INCREMENT.

ER Diagram



ER is the same as the beginning, as explained above.

5. Discuss what functionalities you added or removed. Why?

We added the following functionalities overall to our project. Anything outside the scope of the minimum requirement or any changes from the proposal will be explained in line.

Added functionality based on Proposal

- Inserting Athlete Information (CRUD)
- Deleting Athlete Information (CRUD)
- Searching Athlete Information (CRUD)
- Visualising World Map with Athlete Count
- Stored Procedure + Visualization Bubble chart showing Athlete count based on country/sport
- Filtered view to find athletes based on sport/countries (This was an implementation of an advanced query outside the Stored Procedure - Filtering on the front end)
- Stored Procedure + Filtered view to find athletes based on sport/countries
- The visualization that finds athletes based on sport/countries (This was an implementation of an advanced query outside the Stored Procedure - Filtering on the front end but is done on Python Plotly) - Analytics
- Medal ranking table from a stored procedure

- Medal stats making country ranking visualization (This is a visualization of Trigger)
- Update Medal (CRUD)
- Add/Delete Trigger
- Event Page (Filtered View)

Removed Functionalities

- Based on feasibility & time constraints, the potential to book event spaces (Role/Privilege based) - This was beyond the scope of original requirements and was an add-on, and we wanted to offer functionality within the scope of a conventional Olympic website for now. We had to drop this as it was demanding and we would not be able to complete this on time as it was already beyond the scope during the proposal

6. Explain how you think your advanced database programs complement your application.

Our advanced programs allow us to filter through all of our data. If a user wanted to find athletes from a certain country or sport, they would not have a convenient way to search for that information. Our first advanced query allows us to filter by country and sport to find all the athletes that fit those categories. Our second, advanced query allows us to display countries according to their medal count. This is especially useful in this dataset because each country places importance on its overall placement.

For the Advanced Query 1 + Stored Procedure:

We implemented this functionality in two places. Advanced Query 1 lets me search for Athletes based on Country or Sport Name, and this query forms the backbone of my filtering processes as it lets us easily filter data either from the front or back end.

The AQ1 implemented a standalone API that lets us filter athletes from the front end only, where we will pass over 10000+ rows of data, and we could filter on React directly.

The AQ1 was also implemented in Python to let us visualize the filters directly and as a Stored Procedure where the user is able to select filters and conduct the search through that feature.

For the Advanced Query 2 + Stored Procedure:

We implemented this functionality to let us visualize the athletes and countries vs. sports. This allowed us to visualize the distribution of athletes based on their countries and sports or displayed completely on a map.

For the Trigger

We implemented a way to ensure that the medal count is kept intact, and by doing this, we were able to ensure rankings are correct when updates happen. It is critical that these filters add more functionality as they let us filter effectively, visualize more and keep data integrity in check.

7. Each team member should describe one technical challenge that the team encountered. This should be sufficiently detailed such that another future team could use this as helpful advice if they were to start a similar project or where to maintain your project.

Kavya

A technical challenge that I encountered was the initial setup of the project. The documentation helped us set up the MySQL instance on GCP. However, when we moved on to adding teammates to the project, I was having trouble allowing them to access and edit the SQL instance. To solve this issue, I had to add my team members to the IAM users as editors to grant them full access.

IAM

GRANT ACCESS

REMOVE ACCESS

HELP ASSISTANT

LEARN

PERMISSIONS

RECOMMENDATIONS HISTORY

Permissions for project "My First Project"

These permissions affect this project and all of its resources. [Learn more](#)

☐ Include Google-provided role grants

VIEW BY PRINCIPALS

VIEW BY ROLES

Filter

Enter property name or value

?

|||

Type	Principal	Name	Role	Security insights	Inheritance
	873176782588-compute@developer.gserviceaccount.com	Compute Engine default service account	Editor	6676/6679 excess permissions	
	873176782588@cloudservices.gserviceaccount.com	Google APIs Service Agent	Editor	6679/6679 excess permissions	
	ajitracvicchandran@gmail.com	Ajit Ravichandran	Editor	6482/6685 excess permissions	
	kavyakarthis@gmail.com		Owner	7507/7670 excess permissions	
	sam.kremerman@gmail.com		Editor	6606/6679 excess permissions	
	snehamohan@gmail.com		Editor	6630/6679 excess permissions	

One more issue was that GCP APIs were not working properly through the VM. The VM setup screen API access defaults to minimum API Access. This won't allow the VM to access API methods from backend services that invoke GCP APIs. To connect to the backend SQL server, I had to ensure that the main server had access to this endpoint.

One more issue that we encountered was the billing issues encountered, as we missed terminating the instance and had to add the promo codes from other members to counter the same.

Create an instance

To create a VM instance, select one of the options:

- New VM instance**
Create a single VM instance from scratch
- New VM instance from template
Create a single VM instance from an existing template
- New VM instance from machine image
Create a single VM instance from an existing machine image
- Marketplace
Deploy a ready-to-go solution onto a VM instance

Instance configuration

Name: instance-3
Type: New balanced persistent disk
Size: 10 GB
License type: Free
Image: Debian GNU/Linux 11 (bullseye)

Identity and API access

Service account: Compute Engine default service account

Access scopes: Allow full access to all Cloud APIs

Firewall

Allow HTTP traffic
Allow HTTPS traffic

Pricing summary

Monthly estimate: \$25.46
That's about \$0.03 hourly

Item	Monthly estimate
2 vCPU + 4 GB memory	\$24.46
10 GB balanced persistent disk	\$1.00
Total	\$25.46

In addition, we had faced issues integrating our visualizations to React front end, for which we had to modify the running of Python files to run as a standalone service that was later put onto the React framework.

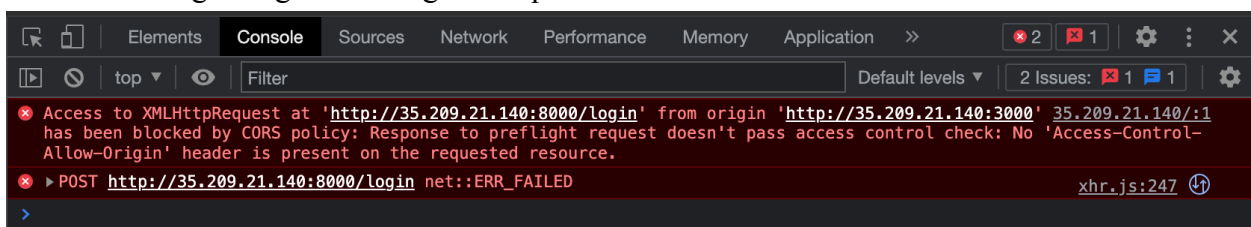
Additionally, the firewall can cause a lot of issues, for which we allowed all IP ranges access across through the firewall over the react port and backend port.

<input type="checkbox"/>	flaskingress	Ingress	Apply to all	IP ranges: 0.0.0.0/0	tcp:3000, 8000 udp:3000, 8000	Allow	▼
--------------------------	------------------------------	---------	--------------	----------------------	----------------------------------	-------	---

Ajit

When designing a backend service, a common problem that anyone will encounter is the CORS issue which would prevent other hosts or ports from utilizing any other port other than the one it is hosted from.

In our project, we host Flask in our backend from port 8000, and when we make requests from any other port other than 8000, we will encounter CORS issues. We have to fix this as the browser level ignoring this setting is complicated and not recommended.



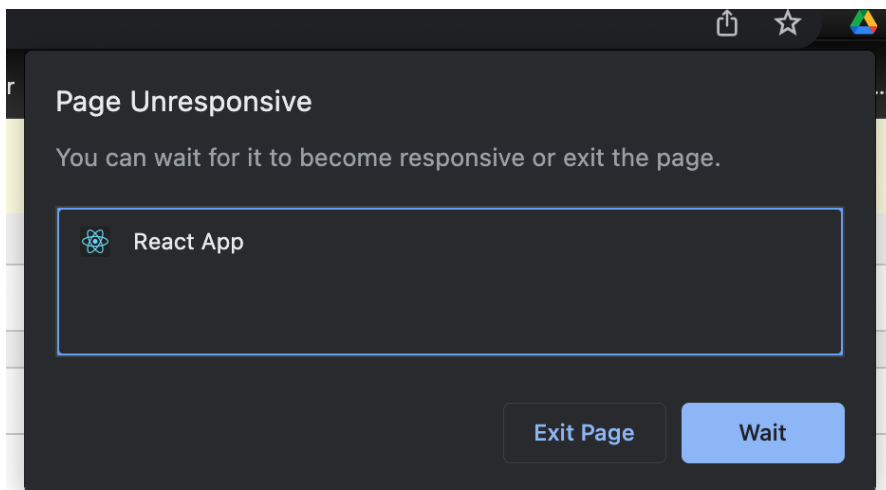
We had to allow CORS login in the backend, especially for POST methods as user login. Updates would not have worked if not, we fixed the issue, which was added in the backend to counter the same.

```
app = Flask(__name__)
CORS(app, resources={r"/": {"origins": ["http://localhost:3000", "http://35.209.21.140:3000"], "methods": ["GET", "POST", "OPTIONS", "DELETE"], "allow_headers": ["Content-Type", "Authorization"]}})
app.config["CORS_HEADERS"] = 'Content-Type'
```

```
@app.route('/login', methods=['POST', 'OPTIONS'])
@cross_origin()
def login():
```

We added the methods above to counter the CORS issue.

Another major issue we ended up encountering was that the backend was returning 10000+ rows in some stages. React had issues rendering the data on screen and straight up conked up the entire website, and we then, as a team, decided to shift to Python for rendering wherever a large number of rows.



Python is very good at handling large numbers of data, and we pass over this react component as an iframe to front end.

```
}
<iframe
  src="http://35.209.21.140:8051"
  title="Dash App"
  style={{ width: "100%", height: "600px", border: "none" }}>
</iframe>
</div>
```

Sneha

When running locally, the code runs on localhost or the loopback IP address 127.0.0.1 or on an internal IP.


```
* Serving Flask app 'backend'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:8000
* Running on http://10.128.0.5:8000
Press CTRL+C to quit
```

While the code can work locally when using localhost, the API would not run on the server. When deployed to the server, the APIs had to be changed everywhere to the server's external IP, as putting localhost did not go well once the server was deployed. Here one can see how the server IP address we have put up to ensure that it can communicate effectively with the backend server.

```
const handleSearch = async () => {
  try {
    const response = await axios.get(`http://35.209.21.140:8000/search_athlete/${searchTerm}`);
    setSearchResults(response.data.search_results);
  }
}
```

We had to test locally with another project copy to ensure the copy runs fine locally and when it is deployed in the server.

```
if __name__ == "__main__":
    app.run_server(debug=True, host="0.0.0.0", port=8051)
```

Additionally, when we implemented the Visualisations, we had to run it on port 0.0.0.0 so that it could accept incoming connections from any port.

Sam

We faced lots of issues handling the dataset provided, our dataset involved was <https://www.kaggle.com/datasets/arjunprasadsarkhel/2021-olympics-in-tokyo> which consists of data that had no connection to one other.

The data looked like this.

Tokyo Data										
File Edit View Insert Format Data Tools Extensions Help										
100% \$ % .00 123 Default... 10 B I A										
A1	A	B	C	D	E	F	G	H	I	J
1										
2										
3	Athletes	Coaches	EntriesGender	Medals	Teams					
4	Name	Name	Discipline/Sport	Rank	Country					
5	Country	Country	Female Count	Country	Discipline/Sport					
6	Discipline/Sport	Discipline/Sport	Male Count	Gold Count	Country					
7		Event	Total	Silver Count	Event					
8				Bronze Count						
9				Total Count						
10				Rank By Total						
11										
12										
13										
14										
15	Login		Country		Sport		Athletes		Coaches	
16	Username		Name		Name		Name		Name	
17	Password		Sport		Medal Count		Male/Female		Country	
18	Name		Athlete		Male/Female		Sport		Discipline/Sport	
19	Email		Medals		Country		Medal Count		Event	
20					Team/Individual		Country		Athlete	
21										
22										
23										
24										
25										
26	Login	ID		Coach	ID					
27		Username			Name					
28		Password			Country					
29		Name			Sport					
30		Email			Event					
31										
32										
33	Athlete	ID		Entries	ID					
34		Country			Sport					
35		Name			Male					
36		Discipline			Female					
37										

As described above, the provided data did not have any direct way to join. We added IDs directly on the CSV file and then added columns.

Athlete			
File Edit View Insert Format Data Tools Extensions Help			
100% \$ % .00 123 Times ... 10 +			
A1	A	B	C
1	AALERUD Katrine		137 395
2	ABAD Nestor		175 395
3	ABAGNALE Giovanni		93 395
4	ABALDE Alberto		175 11
5	ABALDE Tamara		175 11
6	ABALO Luc		68 129
7	ABAROA Cesar		39 395
8	ABASS Abobakr		178 395
9	ABBASALI Hamideh		91 395
10	ABBASOV Islam		13 395
11	ABBINGH Lois		131 199
12	ABBOT Emily		11 395
13	ABBOTT Monica		196 32
14	ABDALLA Abubaker Haydar		151 395

Files				
Athlete	test	Country	Coach	Plays.xlsx
Schedule	users	Medals	Sport	

We used Excel Vlookups and Python panda frames to load data to manipulate and join data which helped us easily do our DDL statements later on.

One more issue was when playing around with CSV data, sometimes the data in the column ended with /r in the CSV, for which we changed DB to strip out /r in the data.

The front and backend had to encounter this proactively and handle this character

```
results = cursor.fetchall()
ranking1 = [
    {"country": result[0].rstrip("\r"), "total": result[1]}
    for result in results
]
```

8. Are there other things that changed comparing the final application with the original proposal?

We made no other changes between our original proposal and our final deliverables.

We were able to deliver all the functionality that we promised initially.

We had to remove one functionality as we initially stated this was an add-on at best, and we had to exclude this as it would've eaten more into our time and would have added more unnecessary effort.

- Based on feasibility & time constraints, the potential to book event spaces (Role/Privilege based) - This was beyond the scope of original requirements and was an add-on, and we wanted to offer functionality within the scope of a conventional Olympic website for now, and we had to drop this as it was demanding and we would not be able to complete this on time as it was already beyond the scope during the proposal

9. Describe future work that you think, other than the interface, that the application can improve on

Other than the interface, one improvement that we could have made was improving interactive components within the application. While we have the functionality to add and delete athletes and update medal counts, we could have found other ways to allow users to interact with the data through creative components. Additionally, we could find additional datasets to allow users to see where specific countries are earning their medals and the outcomes of highly anticipated events. There would have been an additional page that would allow users to scroll through the events for the scores, highlights, and the next time those athletes will compete.

We were bound severely by the dataset as it did not have many fields we would have preferred, such as images and descriptions and many more details from the athlete's perspective, which could have added much more depth to our project.

In the future, we hope to have better datasets that will let us better display more stuff on our App.

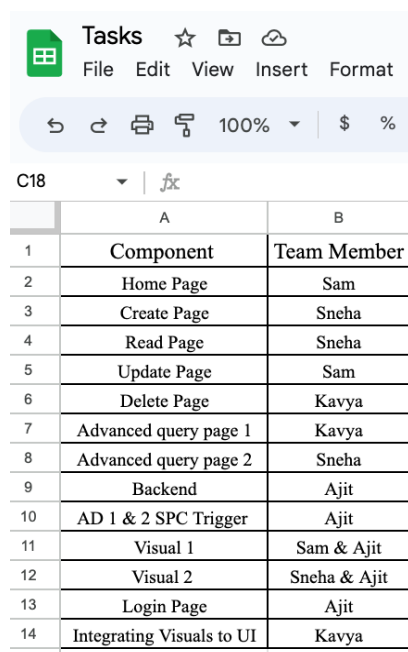
Additionally, we would like to have another backend database like Mongo or similar that can let us add more images which we can add for countries, athletes, etc., as SQL is limited when it comes to adding images. It is not straightforward if not done correctly very prior, and even then, we have to upload all of those images onto GCP and map them effectively, which is a hassle that Mongo could have bypassed.

10. Describe the final division of labor and how well you managed teamwork.

We maintained a spreadsheet containing all the components we envisioned in the proposal stages, and we added or removed functionality based on comforts and strengths.

In terms of the division of labor, we initially sat down in Grainger to review everyone's understanding of concepts since many of the concepts were new to many people on the team. They all sat together, leaned on each other to learn these concepts, and quickly ramped up. Parallely everyone was working on their individual components and frequently discussed over iMessage, Zooms, or in person accordingly. Everyone actively participated in all of the above as needed.

The final spreadsheet is as follows:



The screenshot shows a Google Sheets interface with a spreadsheet titled 'Tasks'. The spreadsheet has two columns: 'Component' and 'Team Member'. The data is as follows:

	A	B
1	Component	Team Member
2	Home Page	Sam
3	Create Page	Sneha
4	Read Page	Sneha
5	Update Page	Sam
6	Delete Page	Kavya
7	Advanced query page 1	Kavya
8	Advanced query page 2	Sneha
9	Backend	Ajit
10	AD 1 & 2 SPC Trigger	Ajit
11	Visual 1	Sam & Ajit
12	Visual 2	Sneha & Ajit
13	Login Page	Ajit
14	Integrating Visuals to UI	Kavya

Sam took care of creating the home and update page and contributed to the creation of the first visual.

Sneha took care of creating the create page, search page, and the second advanced query page and contributed to the creation of the second visual.

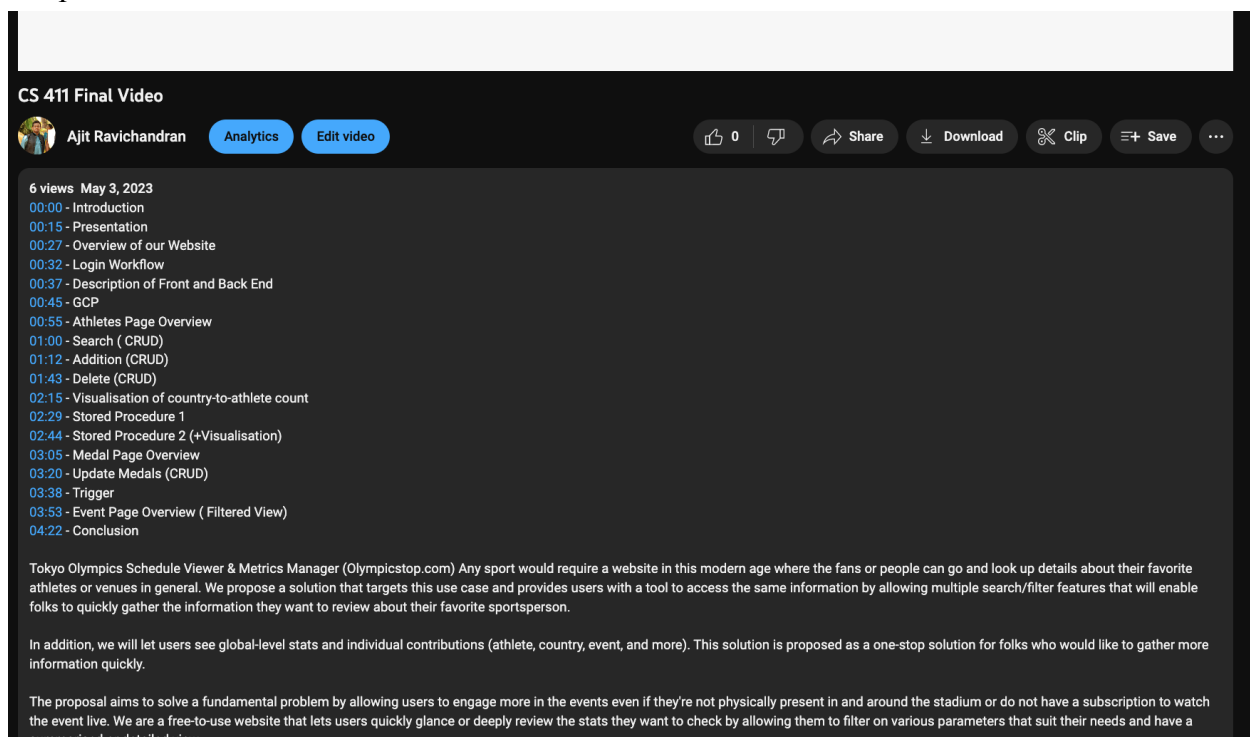
Kavya created the delete page, the first advanced query page, and integrated the visual into the UI.

Ajit took care of the backend alongside all the APIs needed, The Stored Procedures & Triggers for the Advanced Queries, assisting with both the visuals and creating the login page.

While the components were developed in tandem, everyone could push and pull code directly from the server using the stored GitHub SSH keys setup in the server. This ensured that everyone could push and pull code as needed and quickly integrate the code. We coordinated when someone pushed their code, and everyone could work seamlessly without having to fix merge commits that could have emerged from working this way.

Project Demo Video

4:41 seconds - The video comprehensively outlines all the workflows and visualizations and explains the tech stack. The video contains chapter marking points covering all the key components that have to be covered.



Link: <https://youtu.be/aPps5rGYSyQ>